

```
import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

df=pd.read_csv('https://raw.githubusercontent.com/YBI-Foundation/Dataset/main/Bank%20Churn')
```

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   CustomerId            10000 non-null  int64  
1   Surname                10000 non-null  object  
2   CreditScore            10000 non-null  int64  
3   Geography              10000 non-null  object  
4   Gender                 10000 non-null  object  
5   Age                   10000 non-null  int64  
6   Tenure                 10000 non-null  int64  
7   Balance                10000 non-null  float64 
8   Num Of Products        10000 non-null  int64  
9   Has Credit Card        10000 non-null  int64  
10  Is Active Member       10000 non-null  int64  
11  Estimated Salary       10000 non-null  float64 
12  Churn                  10000 non-null  int64  
dtypes: float64(2), int64(8), object(3)
memory usage: 1015.8+ KB
```

```
df.head()
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Pr
0	15634602	Hargrave	619	France	Female	42	2	0.00	
1	15647311	Hill	608	Spain	Female	41	1	83807.86	
2	15619304	Onio	502	France	Female	42	8	159660.80	
3	15701354	Boni	699	France	Female	39	1	0.00	

```
df.duplicated('CustomerId').sum()

0
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   CustomerId            10000 non-null  int64
 1   Surname                10000 non-null  object
 2   CreditScore            10000 non-null  int64
 3   Geography              10000 non-null  object
 4   Gender                 10000 non-null  object
 5   Age                    10000 non-null  int64
 6   Tenure                 10000 non-null  int64
 7   Balance                10000 non-null  float64
 8   Num Of Products        10000 non-null  int64
 9   Has Credit Card        10000 non-null  int64
10   Is Active Member       10000 non-null  int64
11   Estimated Salary       10000 non-null  float64
12   Churn                  10000 non-null  int64
dtypes: float64(2), int64(8), object(3)
memory usage: 1015.8+ KB
```

```
df =df.set_index('CustomerId')
```

```
#encoding
```

```
df['Geography'].value_counts()
```

```
France      5014
Germany     2509
Spain       2477
Name: Geography, dtype: int64
```

```
df.replace({'Geography':{'France':2,'Germany':1,'Spain':0}},inplace=True)
```

```
df['Gender'].value_counts()
```

```
Male        5457
Female      4543
Name: Gender, dtype: int64
```

```
df.replace({'Gender':{'Male':0,'Female':1}},inplace=True)
```

```
df['Num Of Products'].value_counts()
```

```
1      5084
2      4590
3       266
4        60
Name: Num Of Products, dtype: int64
```

```
df.replace({'Num Of Products':{'1':0,2:1,3:1,4:1}},inplace=True)
```

```
df['Has Credit Card'].value_counts()
```

```
1    7055
0    2945
Name: Has Credit Card, dtype: int64
```

```
df['Is Active Member'].value_counts()
```

```
1    5151
0    4849
Name: Is Active Member, dtype: int64
```


```
df.loc[(df['Balance']==0), 'Churn'].value_counts()
```

```
0    3117
1     500
Name: Churn, dtype: int64
```

```
df['Zero Balance']=np.where(df['Balance']>0,1,0)
```

```
df['Zero Balance'].hist()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fe7ae2348d0>



```
df.groupby(['Churn', 'Geography']).count()
```

```
#define Label and Features
```

```
df.columns
```

```
Index(['Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure',
       'Balance', 'Num Of Products', 'Has Credit Card', 'Is Active Member',
       'Estimated Salary', 'Churn', 'Zero Balance'],
      dtype='object')
```

```
X=df.drop(['Surname','Churn'],axis=1)
```

```
y=df['Churn']
```

```
X.shape,y.shape
```


```
((10000, 11), (10000,))
```

```
df['Churn'].value_counts()
```

```
0    7963
1    2037
Name: Churn, dtype: int64
```

```
sns.countplot(x='Churn',data=df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fe7ae2404d0>
```



```
X.shape,y.shape
```

```
((10000, 11), (10000,))
```

```
#random under Sanpling
```

```
from imblearn.under_sampling import RandomUnderSampler
```

```
rus=RandomUnderSampler(random_state=2529)
```

```
X_rus,y_rus=rus.fit_resample(X,y)
```

```
X_rus.shape,y_rus.shape,X.shape,y.shape
```

```
((4074, 11), (4074,), (10000, 11), (10000,))
```

```
y.value_counts()
```


```
0    7963
1    2037
Name: Churn, dtype: int64
```

```
y_rus.value_counts()
```

```
0    2037
1    2037
Name: Churn, dtype: int64
```

```
y_rus.plot(kind='hist')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fe7ac00f390>



```
#Random Over sampling
from imblearn.over_sampling import RandomOverSampler
```

```
ros=RandomOverSampler(random_state=2529)
```

```
X_ros,y_ros=ros.fit_resample(X,y)
```

```
X_ros.shape,y_ros.shape,X.shape,y.shape
```

```
((15926, 11), (15926,), (10000, 11), (10000,))
```

```
y.value_counts()
```

```
0    7963
```

```
1    2037
   Name: Churn, dtype: int64
```

```
#Train test splt
from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=2529) #split

X_train_rus,X_test_rus,y_train_rus,y_test_rus=train_test_split(X_rus,y_rus,test_size=0.3,r

X_train_ros,X_test_ros,y_train_ros,y_test_ros=train_test_split(X_ros,y_ros,test_size=0.3,r

#standardize Features
from sklearn.preprocessing import StandardScaler

sc=StandardScaler()

#standardize original dataset

X_train[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_tr

X_test[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_tes

#Standardize random inder sample data

X_train_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(

X_test_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X

#standardize random over sample data

X_train_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(

X_test_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X

#support vector machine classifier
from sklearn.svm import SVC
```

```
svc=SVC()
```

```
svc.fit(X_train,y_train)
```

```
SVC()
```

```
y_pred=svc.predict(X_test)
```

```
#Model accuracy
```

```
from sklearn.metrics import confusion_matrix,classification_report
```

```
confusion_matrix(y_test,y_pred)
```

```
array([[2390,   24],
       [ 459,  127]])
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.99	0.91	2414
1	0.84	0.22	0.34	586
accuracy			0.84	3000
macro avg	0.84	0.60	0.63	3000
weighted avg	0.84	0.84	0.80	3000

```
#hyperparameter Tunning
```

```
from sklearn.model_selection import GridSearchCV
```

```
param_grid={'C':[0.1,1,10],
            'gamma':[1,0.1,0.01],
            'kernel':['rbf'],
            'class_weight':['balanced']}
```

```
grid=GridSearchCV(SVC(),param_grid,refit=True,verbose=2,cv=2)
```

```
grid.fit(X_train,y_train)
```

```
Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.7s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.4s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
```

```

[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.3s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.3s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
GridSearchCV(cv=2, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10], 'class_weight': ['balanced'],
                         'gamma': [1, 0.1, 0.01], 'kernel': ['rbf']},
             verbose=2)

```

```
print(grid.best_estimator_)
```

```
SVC(C=0.1, class_weight='balanced', gamma=1)
```

```
grid_predictions=grid.predict(X_test)
```

```
confusion_matrix(y_test,grid_predictions)
```

```
array([[2016, 398],
       [ 255, 331]])
```

```
print(classification_report(y_test,grid_predictions))
```

	precision	recall	f1-score	support
0	0.89	0.84	0.86	2414
1	0.45	0.56	0.50	586
accuracy			0.78	3000
macro avg	0.67	0.70	0.68	3000
weighted avg	0.80	0.78	0.79	3000

```
#model with random under sampling
```

```
svc_rus=SVC()
```

```
svc_rus.fit(X_train_rus,y_train_rus)
```

```
SVC()
```

```
y_pred_rus=svc_rus.predict(X_test_rus)
```



```
#Model Accuracy
```

```
confusion_matrix(y_test_rus,y_pred_rus)
```

```
array([[457, 170],
       [202, 394]])
```

```
print(classification_report(y_test_rus,y_pred_rus))
```

	precision	recall	f1-score	support
0	0.69	0.73	0.71	627
1	0.70	0.66	0.68	596
accuracy			0.70	1223
macro avg	0.70	0.69	0.70	1223
weighted avg	0.70	0.70	0.70	1223

```
#hyperparameter tuning
```

```
param_grid={'C':[0.1,1,10],
            'gamma':[1,0.1,0.01],
            'kernel':['rbf'],
            'class_weight':['balanced']}
```

```
grid_rus=GridSearchCV(SVC(),param_grid,refit=True,verbose=2,cv=2)
grid_rus.fit(X_train_rus,y_train_rus)
```

```
Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.3s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.3s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
GridSearchCV(cv=2, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10], 'class_weight': ['balanced'],
                          'gamma': [1, 0.1, 0.01], 'kernel': ['rbf']},
             verbose=2)
```

```

print(grid_rus.best_estimator_)

SVC(C=1, class_weight='balanced', gamma=0.1)

grid_predictions_rus=grid_rus.predict(X_test_rus)

confusion_matrix(y_test_rus,grid_predictions_rus)

array([[462, 165],
       [206, 390]])

print(classification_report(y_test_rus,grid_predictions_rus))

```

	precision	recall	f1-score	support
0	0.69	0.74	0.71	627
1	0.70	0.65	0.68	596
accuracy			0.70	1223
macro avg	0.70	0.70	0.70	1223
weighted avg	0.70	0.70	0.70	1223

```

#model with random over sampling
svc_ros=SVC()

svc_ros.fit(X_train_ros,y_train_ros)

SVC()

y_pred_ros=svc_ros.predict(X_test_ros)

#model accuracy

confusion_matrix(y_test_ros,y_pred_ros)

array([[1807, 572],
       [ 744, 1655]])

print(classification_report(y_test_ros,y_pred_ros))

```

	precision	recall	f1-score	support
0	0.71	0.76	0.73	2379
1	0.74	0.69	0.72	2399
accuracy			0.72	4778
macro avg	0.73	0.72	0.72	4778

weighted avg	0.73	0.72	0.72	4778
--------------	------	------	------	------

```
#hyperparameter tuning
param_grid={'C':[0.1,1,10],
            'gamma':[1,0.1,0.01],
            'kernel':['rbf'],
            'class_weight':['balanced']}
```

```
grid_ros=GridSearchCV(SVC(),param_grid,refit=True,verbose=2,cv=2)
grid_ros.fit(X_train_ros,y_train_ros)
```

```
Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 4.3s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 4.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.6s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.5s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.6s
[CV] END ....C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.5s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.9s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.9s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.0s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.5s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.5s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 3.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.1s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.0s
GridSearchCV(cv=2, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10], 'class_weight': ['balanced'],
                          'gamma': [1, 0.1, 0.01], 'kernel': ['rbf']},
             verbose=2)
```

```
print(grid_ros.best_estimator_)
```

```
SVC(C=10, class_weight='balanced', gamma=1)
```

```
grid_predictions_ros=grid_ros.predict(X_test_ros)
```

```
confusion_matrix(y_test_ros,grid_predictions_ros)
```

```
array([[1990, 389],
       [ 92, 2307]])
```

```
print(classification_report(y_test_ros,grid_predictions_ros))
```

	precision	recall	f1-score	support
0	0.96	0.84	0.89	2379
1	0.86	0.96	0.91	2399
accuracy			0.90	4778

macro avg	0.91	0.90	0.90	4778
weighted avg	0.91	0.90	0.90	4778

✓ 0s completed at 12:15 AM

×