

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt


import seaborn as sns

df=pd.read_csv('https://raw.githubusercontent.com/YBI-Foundation/Dataset/main/MPG.csv')
```

df

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
0	18.0	8	307.0	130.0	3504	12.0	70
1	15.0	8	350.0	165.0	3693	11.5	70
2	18.0	8	318.0	150.0	3436	11.0	70
3	16.0	8	304.0	150.0	3433	12.0	70
4	17.0	8	302.0	140.0	3449	10.5	70
...	...	...	...	...	...	...	...
393	27.0	4	140.0	86.0	2790	15.6	82

df.head()

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	or
0	18.0	8	307.0	130.0	3504	12.0	70	
1	15.0	8	350.0	165.0	3693	11.5	70	

```
df.nunique()
```

```
mpg          129
cylinders     5
displacement  82
horsepower    93
weight       351
acceleration  95
model_year    13
origin        3
name         305
dtype: int64
```

```
#data Preprocessing
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg             398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       392 non-null   float64
4   weight           398 non-null   int64
5   acceleration     398 non-null   float64
6   model_year       398 non-null   int64
7   origin           398 non-null   object
8   name             398 non-null   object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
df.describe() #summary statistics
```

	mpg	cylinders	displacement	horsepower	weight	acceleration
<b>count</b>	398.000000	398.000000	398.000000	392.000000	398.000000	398.000000
<b>mean</b>	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090
<b>std</b>	7.815984	1.701004	104.269838	38.491160	846.841774	2.757689
<b>min</b>	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000
<b>25%</b>	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000
<b>50%</b>	23.000000	4.000000	148.500000	93.500000	2803.500000	15.500000
<b>75%</b>	29.000000	8.000000	262.000000	126.000000	3608.000000	17.175000
<b>max</b>	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000

```
df.corr()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration
<b>mpg</b>	1.000000	-0.775396	-0.804203	-0.778427	-0.831741	0.420289
<b>cylinders</b>	-0.775396	1.000000	0.950721	0.842983	0.896017	-0.505419
<b>displacement</b>	-0.804203	0.950721	1.000000	0.897257	0.932824	-0.543684
<b>horsepower</b>	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196
<b>weight</b>	-0.831741	0.896017	0.932824	0.864538	1.000000	-0.417457
<b>acceleration</b>	0.420289	-0.505419	-0.543684	-0.689196	-0.417457	1.000000
<b>model_year</b>	0.579267	-0.348746	-0.370164	-0.416361	-0.306564	0.288111

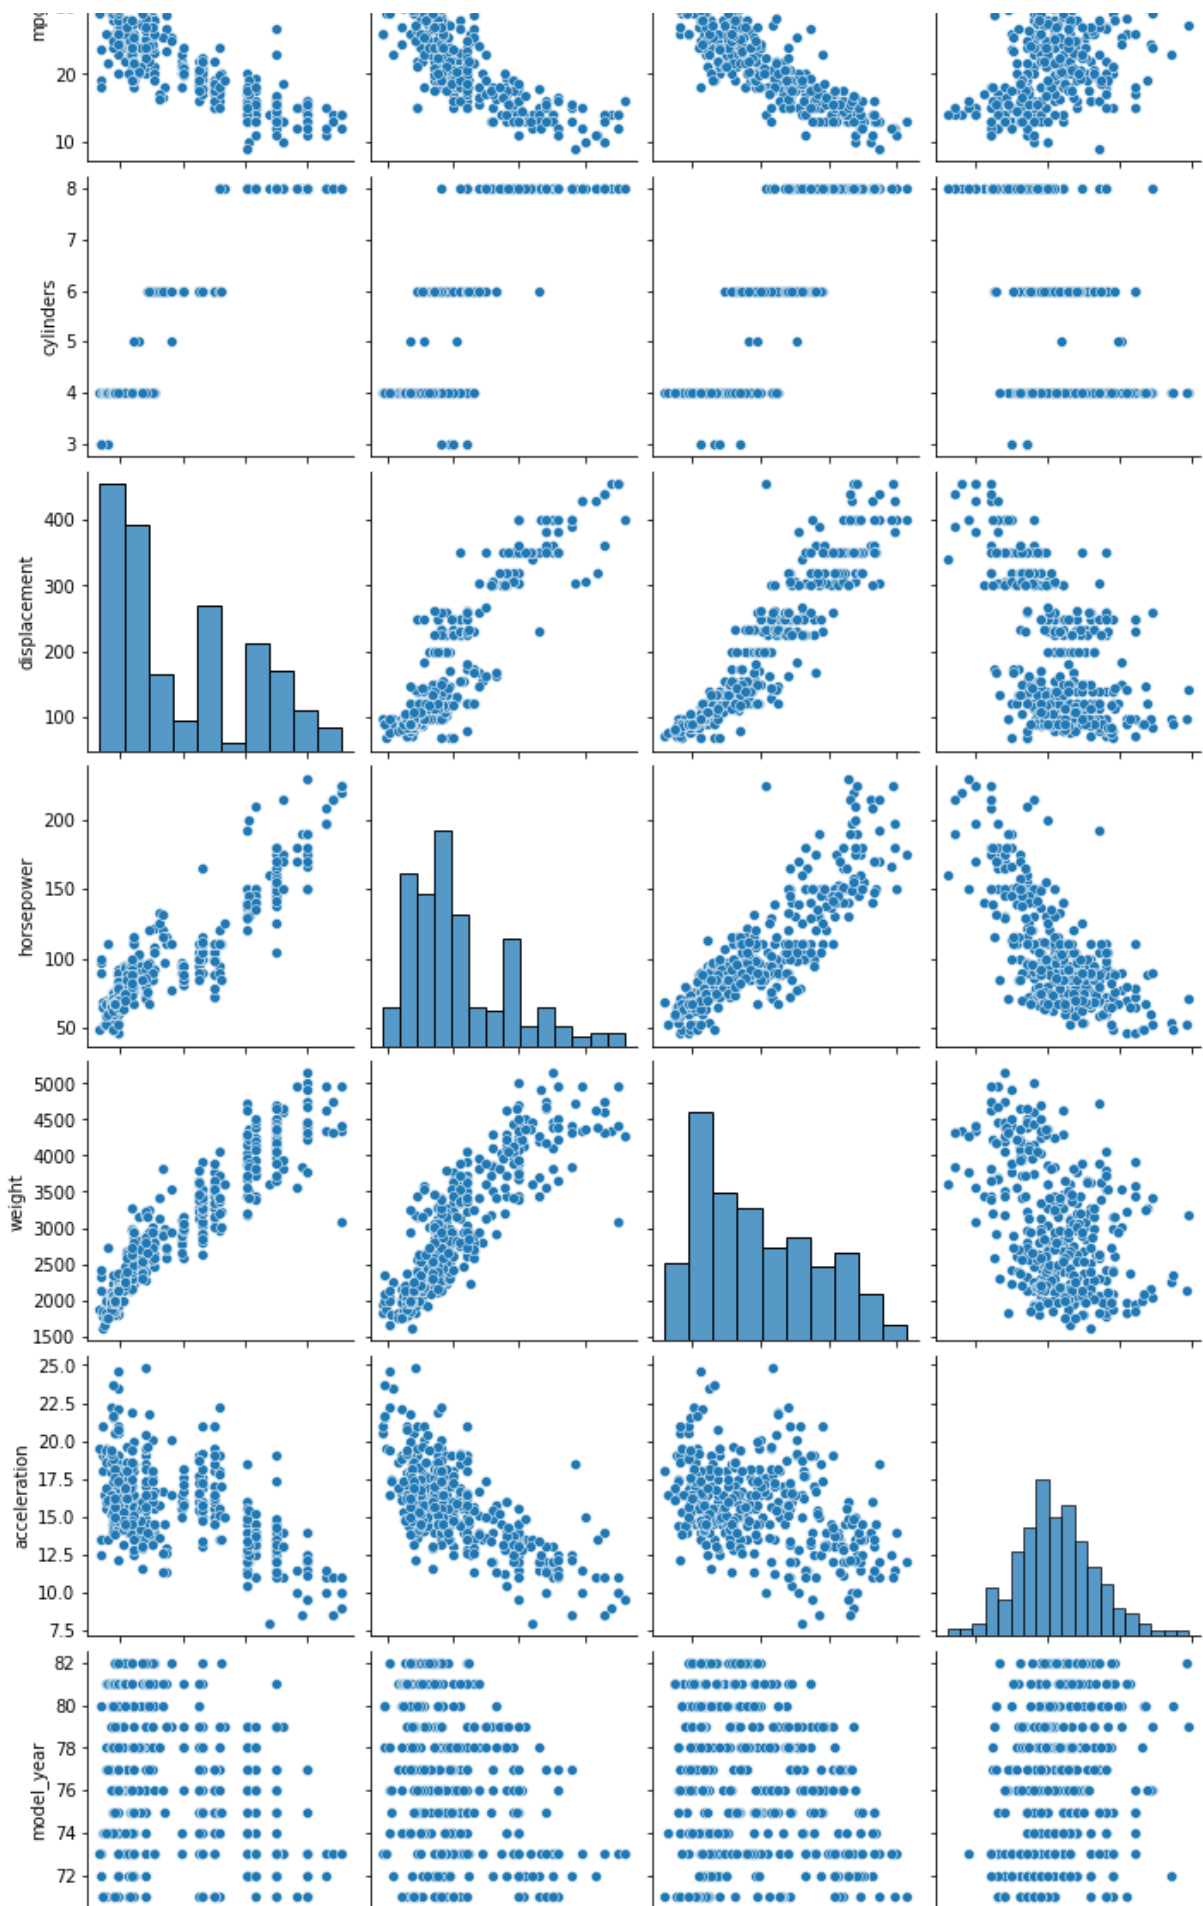
```
df=df.dropna() #get remove missing value
```

```
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null    float64
1   cylinders        398 non-null    int64
2   displacement     398 non-null    float64
3   horsepower       392 non-null    float64
4   weight          398 non-null    int64
5   acceleration     398 non-null    float64
6   model_year      398 non-null    int64
7   origin          398 non-null    object
8   name            398 non-null    object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
#data visualization
```

```
sns.pairplot(df,x_vars=['displacement','horsepower','weight','acceleration'])
```



```
sns.regplot(x='displacement',y='mpg',data=df);
```



```
df.columns
```

```
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
      'acceleration', 'model_year', 'origin', 'name'],
      dtype='object')
```

```
y=df['mpg']
```

```
y.shape
```

```
(392,)
```

```
X=df[['displacement','horsepower','weight','acceleration']]
```

```
X.shape
```

```
(392, 4)
```

```
#scaling data
```

```
from sklearn.preprocessing import StandardScaler
```

```
ss=StandardScaler()
```

```
X=ss.fit_transform(X)
```

```
X
```

```
array([[ 1.07728956,  0.66413273,  0.62054034, -1.285258  ],
       [ 1.48873169,  1.57459447,  0.84333403, -1.46672362],
       [ 1.1825422 ,  1.18439658,  0.54038176, -1.64818924],
       ...,
       [-0.56847897, -0.53247413, -0.80463202, -1.4304305 ],
       [-0.7120053 , -0.66254009, -0.41562716,  1.11008813],
       [-0.72157372, -0.58450051, -0.30364091,  1.40043312]])
```

```
pd.DataFrame(X).describe()
```

	0	1	2	3
<b>count</b>	3.920000e+02	3.920000e+02	3.920000e+02	3.920000e+02
<b>mean</b>	-2.537653e-16	-4.392745e-16	5.607759e-17	6.117555e-16
<b>std</b>	1.001278e+00	1.001278e+00	1.001278e+00	1.001278e+00
<b>min</b>	-1.209563e+00	-1.520975e+00	-1.608575e+00	-2.736983e+00
<b>25%</b>	-8.555316e-01	-7.665929e-01	-8.868535e-01	-6.410551e-01
<b>50%</b>	-4.153842e-01	-2.853488e-01	-2.052109e-01	-1.499869e-02
<b>75%</b>	7.782764e-01	5.600800e-01	7.510927e-01	5.384714e-01
<b>max</b>	2.493416e+00	3.265452e+00	2.549061e+00	3.360262e+00



```
#TRAIN TEST SPLIT DATA
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,train_size=0.7,random_state=2529)
```

```
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
((274, 4), (118, 4), (274,), (118,))
```

```
#linear Regression Model
```

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()
```

```
lr.fit(X_train,y_train)
```

```
LinearRegression()
```

```
lr.intercept_
```

```
23.485738559737584
```

```
lr.coef_
```

```
array([-1.05767743, -1.68734727, -4.10787617, -0.11495177])
```

```
#predict Test data
```

```
y_pred=lr.predict(X_test)
```

```
array([18.51865637, 15.09305675, 14.30128789, 23.6753321 , 29.7546115 ,
       23.68796629, 26.61066644, 24.56692437, 15.06260986, 11.94312046,
       24.08050053, 27.96518468, 31.66130278, 31.01309132, 18.32428976,
       19.32795009, 28.08847536, 32.1506879 , 31.15859692, 27.15792144,
       18.82433097, 22.54580176, 26.15598115, 32.36393869, 20.74377679,
        8.78027518, 22.19699435, 18.20614294, 25.00052718, 15.26421552,
       23.13441082, 17.10542257,  9.87180062, 30.00790415, 20.41204655,
       29.11860245, 24.4305187 , 21.72601835, 10.51174626, 13.12426391,
       21.41938406, 19.96113872,  6.19146626, 17.79025345, 22.5493033 ,
       29.34765021, 13.4861847 , 25.88852083, 29.40406946, 22.41841964,
       22.07684766, 16.46575802, 24.06290693, 30.12890046, 10.11318121,
        9.85011438, 28.07543852, 23.41426617, 20.08501128, 30.68234133,
       20.92026393, 26.78370281, 22.9078744 , 14.15936872, 24.6439883 ,
       26.95515832, 15.25709393, 24.11272087, 30.80980589, 14.9770217 ,
       27.67836372, 24.2372919 , 10.92177228, 30.22858779, 30.88687365,
       27.33992044, 31.18447082, 10.8873597 , 27.63510608, 16.49231363,
       25.63229888, 29.49776285, 14.90393439, 32.78670687, 30.37325244,
       30.9262743 , 14.71702373, 27.09633246, 26.69933806, 29.06424799,
       32.45810182, 29.44846898, 31.61239999, 31.57891837, 21.46542321,
       31.76739191, 26.28605476, 28.96419915, 31.09628395, 24.80549594,
       18.76490961, 23.28043777, 23.04466919, 22.14143162, 15.95854367,
       28.62870918, 25.58809869, 11.4040908 , 25.73334842, 30.83500051,
       21.94176255, 15.34532941, 30.37399213, 28.7620624 , 29.3639931 ,
       29.10476703, 20.44662365, 28.11466839])
```

#Model Accuracy

```
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error,r2_score
```

```
mean_absolute_error(y_test,y_pred)
```

```
3.3286968643244106
```

```
mean_absolute_percentage_error(y_test,y_pred)
```

```
0.14713035779536746
```

```
r2_score(y_test,y_pred)
```

```
0.7031250746717692
```

#Polynomial Regression

```
from sklearn .preprocessing import PolynomialFeatures
```

```
poly=PolynomialFeatures(degree=2,interaction_only=True,include_bias=False)
```

```
X_train2=poly.fit_transform(X_train)
```

```
X_test2=poly.fit_transform(X_test)
```

```
lr.fit(X_train2,y_train)

LinearRegression()

lr.intercept_

21.27336450063766

lr.coef_

array([-2.76070596, -5.00559628, -1.36884133, -0.81225214,  1.24596571,
        -0.12475017, -0.90542822,  1.35064048, -0.17337823,  1.41680398])

y_pred_poly=lr.predict(X_test2)

#Model accuracy
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error,r2_score

mean_absolute_error(y_test,y_pred_poly)

2.7887147720295977

mean_absolute_percentage_error(y_test,y_pred_poly)

0.1207401834293869

r2_score(y_test,y_pred_poly)

0.7461731314563803

#Hand written digit prediction-classification model
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt


#import data
from sklearn.datasets import load_digits

df=load_digits()

_, axes=plt.subplots(nrows=1,ncols=4,figsize=(10,3))
for ax,image,label in zip(axes,df.images,df.target):
    ax.set_axis_off()
    ax.imshow(image,cmap=plt.cm.gray_r,interpolation="nearest")
```



```
ax.set_title("Training:%i" % label)
```



```
#data processing
#Flatten image
df.images.shape
```

```
(1797, 8, 8)
```

```
df.images[0]
```

```
array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.],
       [ 0.,  0., 13., 15., 10., 15.,  5.,  0.],
       [ 0.,  3., 15.,  2.,  0., 11.,  8.,  0.],
       [ 0.,  4., 12.,  0.,  0.,  8.,  8.,  0.],
       [ 0.,  5.,  8.,  0.,  0.,  9.,  8.,  0.],
       [ 0.,  4., 11.,  0.,  1., 12.,  7.,  0.],
       [ 0.,  2., 14.,  5., 10., 12.,  0.,  0.],
       [ 0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
```

```
df.images[0].shape
```

```
(8, 8)
```

```
len(df.images)
```

```
1797
```

```
n_samples=len(df.images)
data=df.images.reshape((n_samples,-1))
```

```
data[0]
```

```
array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
        15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
        12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
         0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
        10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
```

```
data[0].shape
```

```
(64,)
```

```
data.shape
```

```
(1797, 64)
```

```
#Scaling data
```

```
data.min()
```

```
0.0
```

```
data.max()
```

```
16.0
```

```
data=data/16
```

```
data.min()
```

```
0.0
```

```
data.max()
```

```
1.0
```

```
#Train test split data
```

```
from sklearn .model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(data,df.target,test_size=0.7,random_state=2
```

```
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
((539, 64), (1258, 64), (539,), (1258,))
```

```
#Random forest model
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf=RandomForestClassifier()
```

```
rf.fit(X_train,y_train)
```

```
RandomForestClassifier()
```

```
#predict test data
y_pred=rf.predict(X_test)
```

```
y_pred
```

```
array([1, 9, 4, ..., 0, 0, 7])
```

```
#Model Accuracy
```

```
from sklearn.metrics import confusion_matrix,classification_report
```

```
confusion_matrix(y_test,y_pred)
```

```
array([[119,  0,  0,  0,  1,  0,  0,  0,  0,  0],
       [ 0, 130,  1,  1,  0,  2,  0,  0,  0,  0],
       [ 0,  0, 118,  0,  0,  0,  0,  1,  1,  0],
       [ 0,  0,  1, 111,  0,  1,  0,  1,  9,  2],
       [ 0,  3,  0,  0, 115,  0,  0,  2,  2,  0],
       [ 0,  0,  0,  0,  2, 117,  0,  0,  0,  2],
       [ 1,  0,  0,  0,  2,  0, 126,  0,  0,  0],
       [ 0,  0,  0,  0,  1,  0,  0, 130,  0,  0],
       [ 0,  5,  2,  1,  0,  2,  0,  1, 112,  1],
       [ 0,  0,  0,  5,  0,  5,  0,  5,  3, 114]])
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	120
1	0.94	0.97	0.96	134
2	0.97	0.98	0.98	120
3	0.94	0.89	0.91	125
4	0.95	0.94	0.95	122
5	0.92	0.97	0.94	121
6	1.00	0.98	0.99	129
7	0.93	0.99	0.96	131
8	0.88	0.90	0.89	124
9	0.96	0.86	0.91	132
accuracy			0.95	1258
macro avg	0.95	0.95	0.95	1258
weighted avg	0.95	0.95	0.95	1258

---

✓ 0s completed at 12:39 AM

● ×